

ST449 Artificial Intelligence and Deep Learning

Lecture 8

Temporal difference and eligibility traces



Milan Vojnovic

<https://github.com/lse-st449/lectures>

Topics of this lecture

- Temporal-difference (TD) learning
- TD prediction
- Sarsa: on-policy TD control
- Q-learning: off-policy TD control

Time-difference (TD) learning

- **TD learning:** a learning method that combines ideas from Monte Carlo (MC) and Dynamic Programming (DP) methods
- **Similarities with MC methods:** learning directly from experience without a model of the environment
 - Model of the environment: transition probabilities and conditional expected one-step rewards
- **Differences with MC methods:** incremental updates of value function estimates
 - MC methods: **episode-by-episode** incremental updates
 - TD methods: **step-by-step** incremental updates

MC prediction

- A basic **every-visit MC method** for non-stationary environments uses the following update rule for estimation of the value function V^π :

$$V(s_t) \leftarrow V(s_t) + \eta[R_t - V(s_t)]$$

where R_t is the observed return from time t , and η is a constant step size

- This method is referred to as **a constant- η MC method**
- We may regard R_t as a **target**
- The update can be performed only after the return R_t is observed (i.e. after the episode is completed)

TD prediction

- Unlike MC methods, TD methods wait only until next time step
- The simplest TD method (so called TD(0)) is defined by the update:

$$V(s_t) \leftarrow V(s_t) + \eta[r_{t+1} + \gamma V(s_{t+1}) - V(s_t)]$$

- This update rule has $r_{t+1} + \gamma V(s_{t+1})$ as the target
- Considered as a bootstrap method: update in part based on an existing estimate

MC vs TD update

- Note that $V^\pi(s) = \mathbf{E}_\pi[R_t | s_t = s]$ (M)
 $= \mathbf{E}_\pi[\sum_{k=0}^{\infty} \gamma^k r_{t+k+1} | s_t = s]$
 $= \mathbf{E}_\pi[r_{t+1} + \gamma \sum_{k=0}^{\infty} \gamma^k r_{t+k+2} | s_t = s]$
 $= \mathbf{E}_\pi[r_{t+1} + \gamma V^\pi(s_{t+1}) | s_t = s]$ (T)

- MC methods use as the target the random variable in (M)
 - Estimate of unknown expected future returns value
- TD methods use as the target the random variable in (T)
 - Estimate for two reasons: expected value in (T) and estimate of the future returns from the next step
- TD methods can be seen as combining the sample of MC and bootstrapping of DP

TD(0) method for estimating a value function

- **Initialization:** V arbitrary, π policy to be evaluated
- **Repeat** for each episode:
Initialize s

Repeat for each step of the episode:

$a \leftarrow$ action given by π for s

Take action a , observe reward r and next state s'

$$V(s) \leftarrow V(s) + \eta[r + \gamma V(s') - V(s)]$$

$$s \leftarrow s'$$

until s is a terminal state

Backup diagram

- TD methods are referred to as **sample backups**
 - They involve looking ahead to a sample successor state (or state-action pair) using the value of the successor and the reward along the way to compute a backed-up value, and then changing the value of the original state
- MC methods are also sample backups
- DP methods are full backups
- The backup diagram of TD(0):

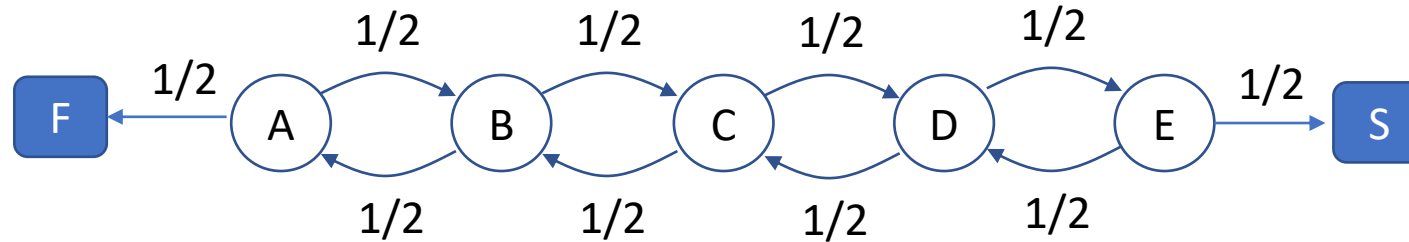


Advantages of TD prediction methods

- Bootstrapping principle: learn a guess from a guess
- Advantage over DP methods: no need to know the environment
- Naturally implementable in a fully incremental fashion
 - Step-by-step as opposed to episode-by-episode increments of MC methods
- TD vs MC with respect to convergence speed:
 - No formal convergence speed theorem
 - In practice, TD methods are usually faster than constant- η MC method

Example: random walk prediction problem


- Consider a simple random walk on a path:



- Reward for transition to state S of value 1, zero reward for any other transition
- Exercise: show that $V(A) = \frac{1}{6}$, $V(B) = \frac{2}{6}$, $V(C) = \frac{3}{6}$, $V(D) = \frac{4}{6}$, $V(E) = \frac{5}{6}$
- Exercise: implement MC and TD methods for estimating $V(s)$ for $s \in S$ (seminar)

Sarsa: an on-policy TD control

- **Sarsa**: a TD prediction methods for the control problem
 - Follows the pattern of generalized policy iteration
 - Uses a TD method for the policy evaluation part
- **Key step**: estimate action value $Q^\pi(s, a)$ for given policy π for all states s and actions a
- An episode interpreted as an alternating sequence of state and state-action pairs:



Sarsa: an on-policy TD control (cont'd)

- The incremental estimation of the action value function:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \eta[r_{t+1} + \gamma Q(s_{t+1}, a_{t+1}) - Q(s_t, a_t)]$$

which is run for any nonterminal state s_t

If s_{t+1} is a terminal state, then $Q(s_{t+1}, a_{t+1}) = 0$

- The update uses every element of the quintuple of variables:

$$(s_t, a_t, r_{t+1}, s_{t+1}, a_{t+1})$$

S A R S A

Sarsa algorithm

- **Initialization:** Q arbitrarily
- **Repeat** for each episode:
 - Initialize s
 - Choose a from s using policy derived from Q (e.g. ϵ -greedy)

Repeat for each step of episode:

Take action a , observe reward r and next state s'

Choose a' from s' using policy derived from Q (e.g. ϵ -greedy)

$$Q(s, a) \leftarrow Q(s, a) + \eta[r + \gamma Q(s', a') - Q(s, a)]$$

$$s \leftarrow s', a \leftarrow a'$$

until s is a terminal state

Q-learning: an off-policy TD control

- Basic one-step Q-learning update:

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \eta[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)]$$

- Learning optimal action value function Q^*
 - For whichever behavior policy has been followed
 - The behavior policy has an effect on which state-action pairs are visited
- Q: Why is Q-learning an off-policy control method?

Q-learning algorithm

- **Initialization:** Q arbitrary
- **Repeat** for each episode:

Initialize s

Repeat for each step of episode:

Choose a from s using policy derived from Q (e.g. ϵ -greedy)

Take action a , observe r and s'

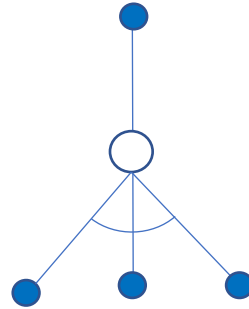
$$Q(s, a) \leftarrow Q(s, a) + \eta \left[r + \gamma \max_{a'} Q(s', a') - Q(s, a) \right]$$

$$s \leftarrow s'$$

until s is a terminal state

Q-learning algorithm (cont'd)

- The backup diagram:



- Q-learning has a guaranteed convergence by the condition that all (state, action) pairs continue to be updated

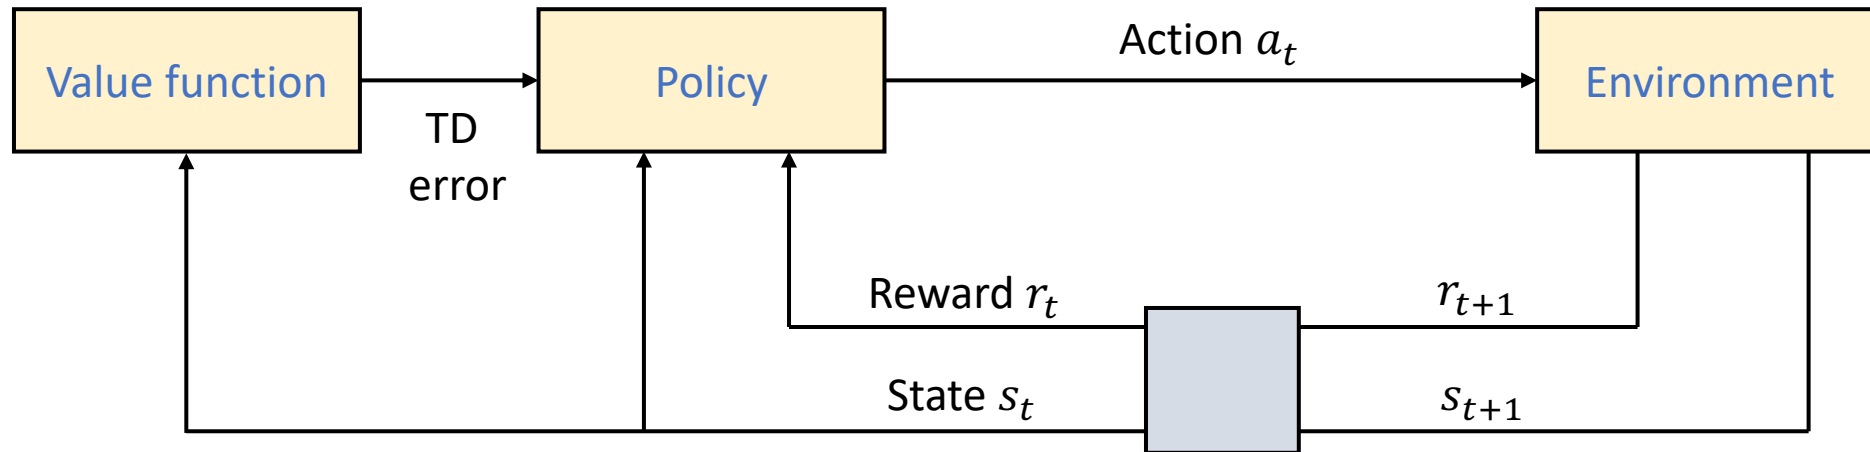
A: why is Q-learning an off-policy learning method?

- Q-learning is an off-policy learning method because it updates Q values using the Q value of the next state s' and action a' selected by using the greedy policy
 - Not using the behavior policy in the action value estimation
- SARSA is an on-policy learning method because it updates Q values using the Q value of next state s' and action a' selected by using the current policy
- Q-learning is an on-policy in the special case when the current policy is a greedy policy, which is of no practical interest because the current policy should explore

Actor-critic methods

- **Actor-critic methods**: TD methods that have a separate memory structure to explicitly represent the policy and value function
 - **Actor**: policy improvement
 - **Critic**: value prediction
- Learning is always on-policy
 - The critic must learn about and critique whichever policy is currently being followed by the actor

The actor-critic architecture



Actor-critic example

- **Critic** maintains value function estimate V and after each action selection, it evaluates the **TD error**:

$$\delta_t = r_{t+1} + \gamma V(s_{t+1}) - V(s_t)$$

- **Actor** uses the TD error for selection of future actions
 - If $\delta_t > 0$ increase preference for selecting action a_t
 - If $\delta_t < 0$ decrease preference for selectin action a_t

- For example, actor may use the policy $\pi_t(s, a) = \frac{e^{p(s,a)}}{\sum_{a'} e^{p(s,a')}}$

where $p(s_t, a_t) \leftarrow p(s_t, a_t) + \beta \delta_t$

An alternative: $p(s_t, a_t) \leftarrow p(s_t, a_t) + \beta (1 - \pi_t(s_t, a_t)) \delta_t$

Eligibility traces

- **Eligibility traces**: a bridge between TD and MC methods
- **(Theoretical) forward view**: eligibility traces produce a family of methods that span a spectrum of methods with MC at one end and one-step TD at the other
- **(Mechanistic) backward view**: an eligibility trace is a temporary record of the occurrence of an event, such as visiting of a state or taking an action
 - The trace marks the memory parameters associated with the event as eligible or undergoing learning changes
- Eligibility traces can be seen as a mechanism for **temporal credit assignment**
 - Similar to TD methods: when a TD error occurs, only the eligible states or actions are assigned credit or blame for the error

n-step TD prediction

- **Goal**: estimate state value function V^π from sample episodes generated by following policy π
- Backup methods:
 - **Full backup**: perform a backup for each state based on **entire sequence of observed rewards from that state until the end of episode** (ex. MC method)
 - **1-step backup**: perform a backup based on **just one next reward**, using the value of the state one step later as a proxy for the remaining rewards (ex. TD(0) method)
 - **n-step backup**: perform a backup based on **next n rewards** and using the value of the state after n steps as a proxy for the remaining rewards (n-step TD methods)

n-step TD prediction cont'd

- Consider the backup applied to state s_t as a result of a state-reward sequence $s_t, r_{t+1}, s_{t+1}, r_{t+1}, \dots, r_T, s_T$ where T is the last step of the episode
- Backup targets:
 - Full: $R_t = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{T-t-1} r_T$
 - 1-step: $R_t^{(1)} = r_{t+1} + \gamma V_t(s_{t+1})$
 - 2-step: $R_t^{(2)} = r_{t+1} + \gamma r_{t+2} + \gamma^2 V_t(s_{t+2})$
 - \vdots
 - n-step: $R_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V_t(s_{t+n})$ (n-step return)
- Truncation: if $n \geq T - t$, then $R_t^{(n)} = R_t^{(T-t)} \equiv R_t$

n-step backup

- The n-step backup is defined by

$$V(s_t) \leftarrow V(s_t) + \Delta V_t(s_t)$$

where

$$\Delta V_t(s_t) = \eta \left(R_t^{(n)} - V_t(s_t) \right)$$

η is a positive step-size parameter, and

$$R_t^{(n)} = \begin{cases} r_{t+1} + \gamma r_{t+2} + \cdots + \gamma^{n-1} r_{t+n} + \gamma^n V_t(s_{t+n}) & \text{if } n < T - t \\ R_t & \text{otherwise} \end{cases}$$

Online vs offline updates

- **Online updates:** updates are done per each time step during the episode **as soon as the increment is computed**

$$V_{t+1}(s) = V_t(s) + \Delta V_t(s) \text{ for } s \in S$$

- **Offline updates:** value function increments are accumulated per step during an episode but are used for updating **only at the end of the episode**

$$V(s) \leftarrow V(s) + \sum_{t=0}^{T-1} \Delta V_t(s) \text{ for } s \in S$$

The forward view and complex backups

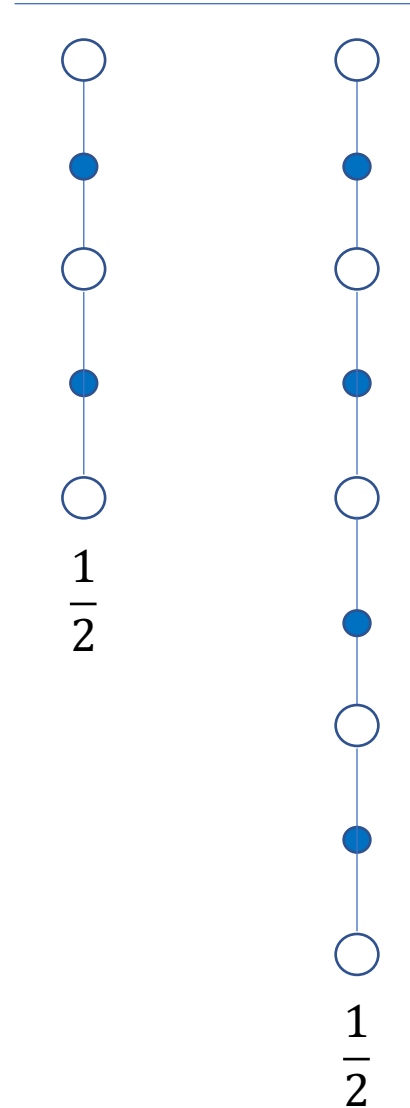
- Backups can be performed not just towards any n-step return, but towards any average of n-step returns
- E.g. average of 2-step and 4-step returns:

$$R_t^{avg} = \frac{1}{2} R_t^{(2)} + \frac{1}{2} R_t^{(4)}$$

- **Complex backups**: a backup that averages simpler component backups

Backup diagram of a complex backup

- Mixing half of a 2-step backup and half of a 4-step backup



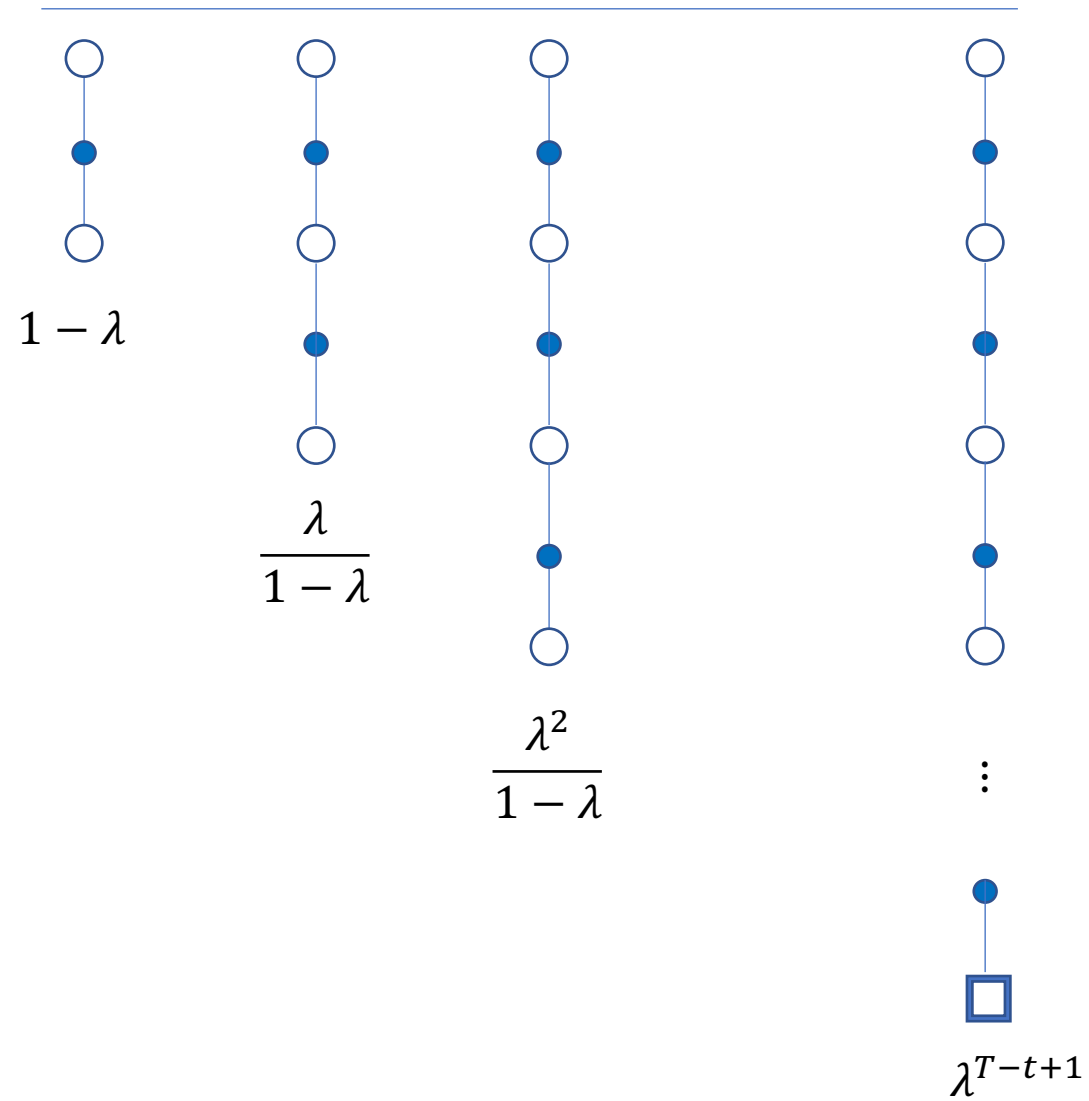
TD(λ) method

- TD(λ) method: uses an average of n -step backups, each weighted with λ^n where $\lambda \in [0,1]$ is a parameter
- The λ -return:

$$R_t^\lambda = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_t^{(n)}$$

- Special cases:
 - Case $\lambda = 0$: $R_t^\lambda = R_t^{(1)}$ (1-step TD method)
 - Case $\lambda = 1$: $R_t^\lambda = R_t$ (constant- η MC method)

Backup diagram for TD(λ)



λ -return algorithm

- For each time step t :

$$\Delta V_t(s) = \begin{cases} \eta [\textcolor{green}{R}_t^\lambda - \textcolor{red}{V}_t(s_t)] & \text{if } s = s_t \\ 0 & \text{otherwise} \end{cases}$$

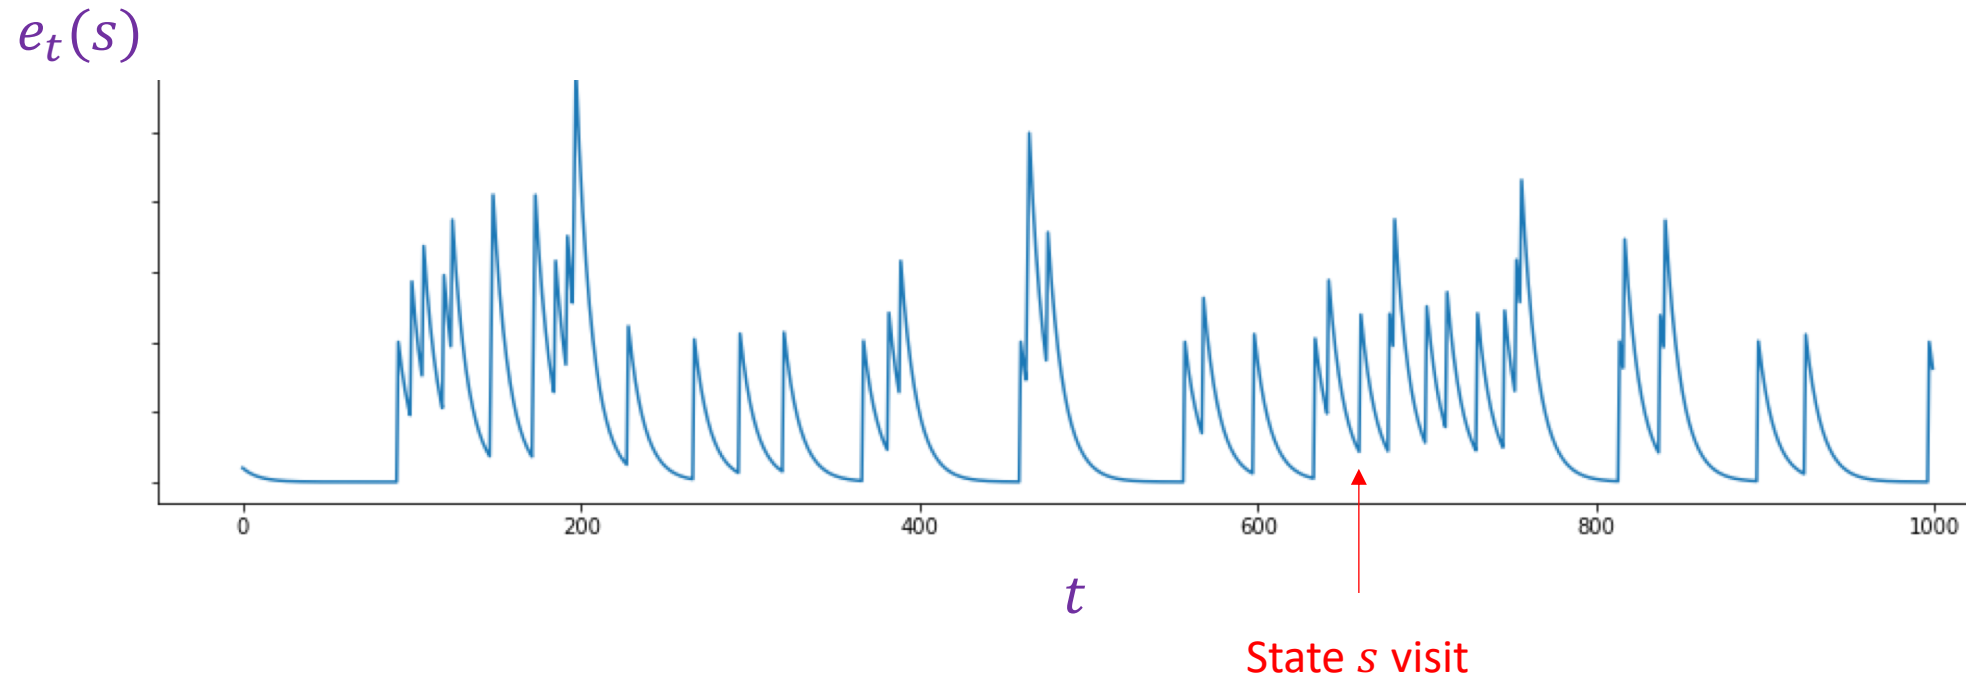
The backward view of TD(λ)

- **Backward view:** causal incremental method for approximating the forward view
- **Eligibility trace:** for state $s \in S$, eligibility trace $e_t(s) \in \mathbf{R}^+$ has jumps of size 1 at each state s visit and exponential time discounting:

$$e_t(s) = \gamma\lambda e_{t-1}(s) + \mathbf{1}_{\{s_t=s\}}$$

- Each eligibility trace indicates the degree to which the corresponding state is eligible for undergoing learning changes

Eligibility trace illustration



On-line TD(λ) algorithm

- Initialization: $V(s)$ arbitrary and $e(s) = 0$ for all $s \in S$
- **Repeat** for each episode:
 - Initialize s
 - Repeat for each step of the episode:
 - $a \leftarrow$ action given by policy π for s
 - Take action a , observe reward r and next state s'
 - $\delta \leftarrow r + \gamma V(s') - V(s)$
 - $e(s) \leftarrow e(s) + 1$
 - For each** s'' :
 - $V(s'') \leftarrow V(s'') + \eta e(s'') \delta$
 - $e(s'') \leftarrow \gamma \lambda e(s'')$
 - $s \leftarrow s'$
 - until** s is terminal

Equivalence of forward and backward views

- Forward view update: $\Delta V_t^\lambda(s_t) = \eta \left(R_t^\lambda - V_t(s_t) \right)$
- Backward view update: $\Delta V_t^{TD}(s) = \eta e_t(s) \delta_t$
- The equivalence property: for offline updating it holds

$$\sum_{t=0}^{T-1} \Delta V_t^{TD}(s) = \sum_{t=0}^{T-1} \Delta V_t^\lambda(s_t) \mathbf{1}_{\{s_t=s\}} \text{ for all } s \in S$$

- For online updating, the equality holds approximately for small step size η (small changes of the value function within an episode)

Proof of the equivalence property

- Basic fact:

$$e_t(s) = \sum_{k=0}^t (\gamma\lambda)^{t-k} \mathbf{1}_{\{s_k=s\}}$$

- $$\begin{aligned} \sum_{t=0}^{T-1} \Delta V_t^{TD}(s) &= \sum_{t=0}^{T-1} \eta \delta_t e_t(s) \\ &= \sum_{t=0}^{T-1} \eta \delta_t \sum_{k=0}^t (\gamma\lambda)^{t-k} \mathbf{1}_{\{s_k=s\}} \\ &= \sum_{k=0}^{T-1} \eta \sum_{t=0}^k (\gamma\lambda)^{k-t} \delta_k \mathbf{1}_{\{s_t=s\}} \\ &= \sum_{t=0}^{T-1} \eta \mathbf{1}_{\{s_t=s\}} \sum_{k=t}^{T-1} (\gamma\lambda)^{k-t} \delta_k \end{aligned}$$

Proof of the equivalence property (cont'd)

- Basic facts: (a) $\frac{1}{\eta} \Delta V_t^\lambda(s_t) = R_t^\lambda - V_t(s_t)$
(b) $R_t^{(n)} = r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n V_t(s_{t+n})$
- $R_t^\lambda - V_t(s_t) = (1 - \lambda) \sum_{n=1}^{\infty} \lambda^{n-1} R_t^{(n)} - V_t(s_t)$
 $= (1 - \lambda) \sum_{n=0}^{\infty} \lambda^n [\sum_{k=0}^n \gamma^k r_{t+k+1} + \gamma^{n+1} V_t(s_{t+n+1})] - V_t(s_t)$
 $= \sum_{n=0}^{\infty} (\gamma \lambda)^n r_{t+n+1} + \sum_{n=0}^{\infty} [(\gamma \lambda)^n \gamma V_t(s_{t+n+1}) - (\gamma \lambda)^{n+1} V_t(s_{t+n+1})] - V_t(s_t)$
 $= \sum_{n=0}^{\infty} (\gamma \lambda)^n r_{t+n+1} + \sum_{n=0}^{\infty} (\gamma \lambda)^n [\gamma V_t(s_{t+n+1}) - V_t(s_{t+n})]$
 $= \sum_{n=0}^{\infty} (\gamma \lambda)^n [r_{t+n+1} + \gamma V_t(s_{t+n+1}) - V_t(s_{t+n})]$
 $= \sum_{k=t}^{\infty} (\gamma \lambda)^{k-t} \delta_k = \sum_{k=t}^{T-1} (\gamma \lambda)^{k-t} \delta_k$

Proof of the equivalence property (cont'd)

- From the last slide, it follows

$$\begin{aligned}\sum_{t=0}^{T-1} \Delta V_t^\lambda(s_t) \mathbf{1}_{\{s_t=s\}} &= \sum_{t=0}^{T-1} \eta \left(\sum_{k=t}^{T-1} (\gamma\lambda)^{k-t} \delta_k \right) \mathbf{1}_{\{s_t=s\}} \\ &= \sum_{t=0}^{T-1} \eta \mathbf{1}_{\{s_t=s\}} \sum_{k=t}^{T-1} (\gamma\lambda)^{k-t} \delta_k\end{aligned}$$

- The proof follows because we have shown that $\sum_{t=0}^{T-1} \Delta V_t^{TD}(s)$ is also equal to the right-hand side in the above equation

Sarsa(λ)

- Eligibility traces are used not just for prediction, but also for control
- A popular approach is to learn the action values $Q_t(s, a)$
- **Basic idea:** apply the TD(λ) prediction method to state-action pairs

$$Q_{t+1}(s, a) = Q_t(s, a) + \eta e_t(s, a) \delta_t \text{ for all } s, a$$

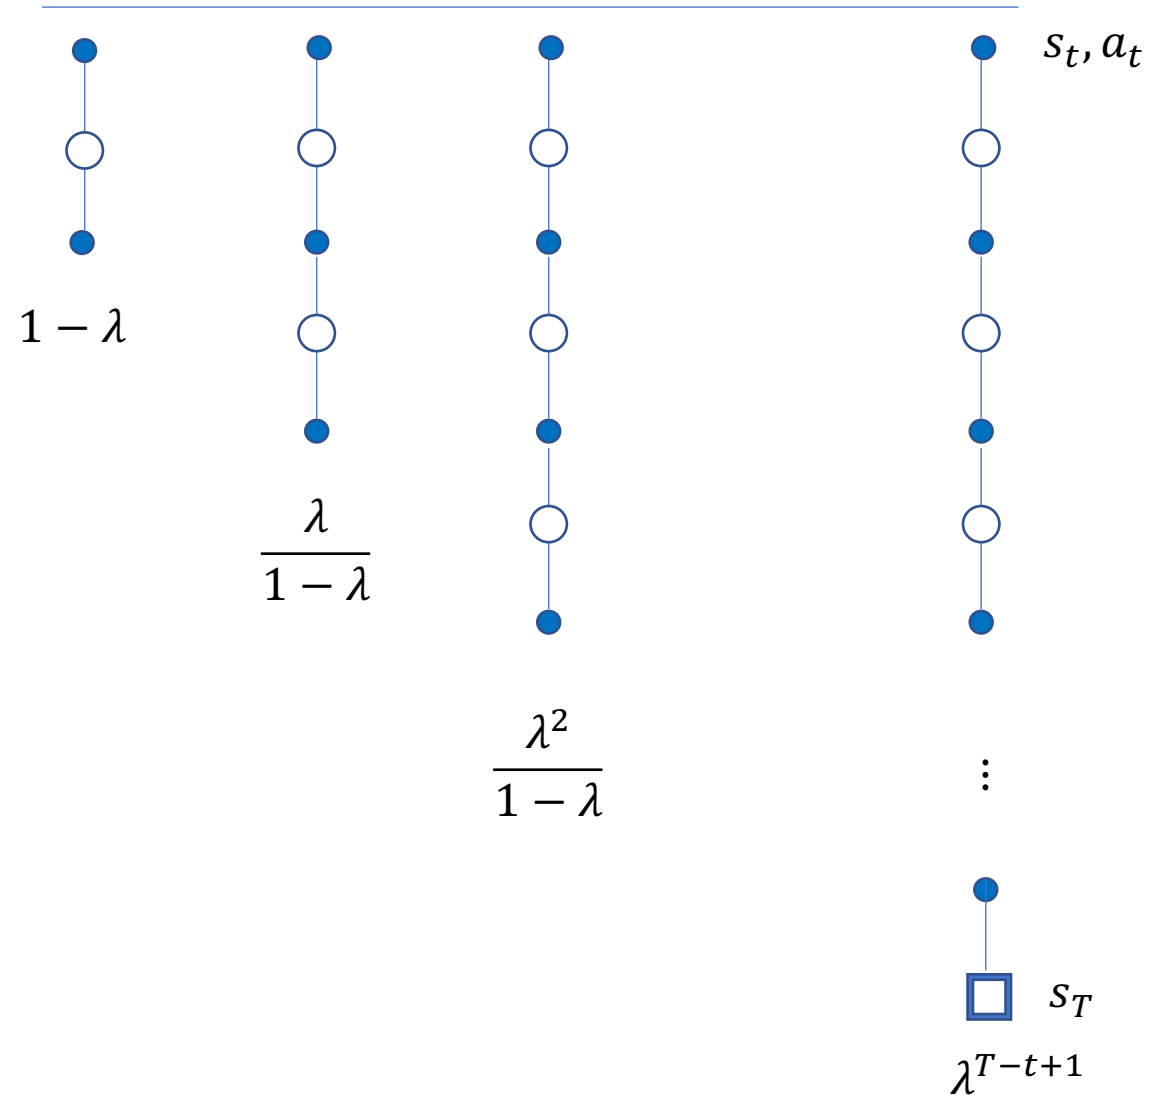
$$\text{where } \delta_t = r_{t+1} + \gamma Q_t(s_{t+1}, a_{t+1}) - Q_t(s_t, a_t)$$

$$\text{and } e_t(s, a) = \gamma \lambda e_{t-1}(s, a) + \mathbf{1}_{\{s=s_t\} \cap \{a_t=a\}}$$

Sarsa(λ) algorithm

- Initialization: $Q(s, a)$ arbitrary and $e(s, a) = 0$ for all s, a
- **Repeat** for each episode:
 - Initialize s, a
 - Repeat** for each step of episode:
 - Take action a , observe r, s'
 - Choose a' from s' using policy derived from Q (ex ϵ -greedy)
 - $\delta \leftarrow r + \gamma Q(s', a') - Q(s, a)$
 - $e(s, a) \leftarrow e(s, a) + 1$
 - For each** s'', a'' :
 - $Q(s'', a'') \leftarrow Q(s'', a'') + \eta e(s'', a'') \delta$
 - $e(s'', a'') \leftarrow \gamma \lambda e(s'', a'')$
 - $s \leftarrow s'$ and $a \leftarrow a'$
 - until** s is terminal

Backup diagram for Sarsa(λ)



References

- R. S. Sutton and A. G. Barto, Reinforcement Learning, Chapters 6 and 7, 1998

Seminar exercises

- TD(0) for random walk problem
- Sarsa: The Windy GridWorld
- Q-learning: The Cliff Walking problem

Extras

MC vs TD(0) exercise: you are the predictor

- Observed (state, reward) outcomes for eight episodes:

1. $(A, 0), (B, 0)$

2. $(B, 1)$

...

7. $(B, 1)$

8. $(B, 0)$

- Q1: What are estimates of $V(A)$ and $V(B)$ according to the MC method?
- Q2: What are estimates of $V(A)$ and $V(B)$ according to the TD(0) method?
- Discuss your answers

The error reduction property

- For any value function V :

$$\max_s |\mathbf{E}_\pi [R_t^{(n)} \mid s_t = s] - V^\pi(s)| \leq \gamma^n \max_s |V(s) - V^\pi(s)|$$

- By using the error reduction property, it can be shown that online and offline TD prediction methods using n -step backups converge to the correct predictions under certain technical conditions

Example: random walk revisited

- Suppose in the first episode the sequence of states is C, D, E
- Assume $V_0(s) = 1/2$ for all $s \in S$
- 1-step method: changes the value estimate only for the last state, $V(E)$, which is incremented towards 1
- 2-step method: changes the value estimates only for the last two states, $V(D)$, $V(E)$, both incremented towards 1
- Three (or more)-n step methods: change the value estimates for all states, $V(C)$, $V(D)$, $V(E)$, incrementing them all towards 1
- Q: which method is better?

Q(λ) learning method

- Q-learning: an off-policy method
 - Learning about the greedy policy while following another policy (typically a policy involving exploratory actions)
 - This requires a care when introducing eligibility traces
- Example:
 - Consider backing up the state-action pair s_t, a_t at time t
 - Suppose in the next two time steps the agent selects greedy actions, but in the on the third it selects an exploratory nongreedy action
 - For learning the value of the greedy policy we can only use the two-step returns
- In general, subsequent experience can be used as long as the greedy policy is being followed

Two $Q(\lambda)$ learning approaches

- **Watkins's $Q(\lambda)$** : only looks ahead as far as the next exploratory action
 - Other principles are much alike to $TD(\lambda)$ and $Sarsa(\lambda)$
 - If a_{t+n} is the first exploratory action, then the longest backup is toward:

$$r_{t+1} + \gamma r_{t+2} + \dots + \gamma^{n-1} r_{t+n} + \gamma^n \max_a Q_t(s_{t+n}, a)$$

- **Peng's $Q(\lambda)$** : avoids cutting off traces every time an exploratory action is taken
 - Cutting off traces every time an exploratory action is taken loses advantage of using eligibility traces
 - The implementation is more complex

Watkins's $Q(\lambda)$ algorithm

- Initialization: $Q(s, a)$ arbitrary and $e(s, a) = 0$ for all s, a
- **Repeat** for each episode:
 - Take action a , observe r, s'
 - Choose a' from s' using policy derived from Q (ex ϵ -greedy)
 - $a^* \leftarrow \operatorname{argmax}_a Q'(s', b)$ (if a' is in argmax then take $a^* \leftarrow a'$)
 - $\delta \leftarrow r + \gamma Q(s', a^*) - Q(s, a)$
 - $e(s, a) \leftarrow e(s, a) + 1$
 - For each** s'', a'' :
 - $Q(s'', a'') \leftarrow Q(s'', a'') + \eta \delta e(s'', a'')$
 - If $a' = a^*$ then $e(s'', a'') \leftarrow \gamma \lambda e(s'', a'')$
 - else $e(s'', a'') \leftarrow 0$
 - $s \leftarrow s'$ and $a \leftarrow a'$
- until** s is terminal

Replacing traces

- A capping of eligibility trace: a replacing trace for state s defined by

$$e_t(s) = \begin{cases} \gamma\lambda e_{t-1}(s) & \text{if } s \neq s_t \\ 1 & \text{if } s = s_t \end{cases}$$

- Different from the standard definition according to which an eligibility trace can assume a value larger than 1
- It can produce a significant improvement in learning rate